



# KubeHound & Beyond:

Evolving security through graphs & automation



Jeremy Fox

March 2024



# Jeremy Fox

Jeremy Fox is a cybersecurity specialist with 10 years experience across government and private sector. Following a career change from the finance industry, he developed a wide range of skills in offensive security from reverse engineering and exploit development to red team operations and cloud security.

He is an engineer at heart, having programmed in everything from C/C++ and ASM, through Python and .NET, to Golang. Although his first love was, and always will be, low-level Windows internals, he now works as a Senior Security Engineer at Datadog developing automated offensive security tooling to detect vulnerabilities in large scale cloud environments. His most recent project is KubeHound, an automated Kubernetes attack path calculator.



# Agenda

---

**01** Introduction

---

**02** The Problem Space

---

**03** The Solution

---

**04** KubeHound In Action

---

**05** Development Process

---

**06** Conclusion

---

**07** Q&A





# Introduction

K8s 101



# Kubernetes 101

## Kubernetes

Open-source container **orchestration platform**

- Automates the deployment, scaling, and management of **containerized applications**
- High availability and auto-scaling

## Container

Lightweight, standalone, and executable software packages

- Encapsulate an application and its dependencies
- **Sandboxed** execution

## Pod

Smallest **deployable unit** in Kubernetes

- Contain one or more containers that **share** the same **network** namespace and **storage** volumes
- Designed to run a **single instance** of an application and are scheduled to *nodes*

## Node

Worker **machines** within a Kubernetes cluster

- **Host pods** and provide the necessary resources (CPU, memory, storage) for running containers
- Grouped together in a **cluster**



# Kubernetes Security 101

## Container escape

Exploit a **container misconfiguration** to gain node access

- Multiple methods of attack
- Very **powerful** - grants access to all node resources

## Kubernetes Identity

Define **service accounts** (robot), user (humans) and groups (both)

- Service accounts define pod identity

## Kubernetes Roles

Set of **permissions** granted to an identity on specific resources

- Addition only (**no deny**)
- Certain permissions are very **powerful** - secrets/list, pods/exec, etc.

## Mounted Volumes

Node or “projected” directories can be mounted into the container

- Mounting the wrong directory = **container escape**
- Projected contain service **account tokens**





# The Problem Space

Scale, complexity and quantifying security

# Vulnerability Context



## FINDING: Overprivileged Container

Web application exposed to the internet running inside a container configured with ***privileged=true***

- Internet facing
- Privilege not required
- Limited auditing



## FINDING: Overprivileged Container

Control plane DNS container running with ***CAP\_SYS\_MODULE*** enabled

- Internal service
- Restricted, audited access
- Privilege required for functionality



# How secure is your cluster?

On a scale of 1-10



# Quantifying Security Posture

**If you cannot measure it, you cannot improve it**

## Current State

*What is the shortest exploitable path between an internet facing service and cluster admin?*

*What percentage of internet-facing services have an exploitable path to cluster admin?*

## Measuring Change

*What type of control would cut off the largest number of attack paths in your cluster?*

*By what percentage did the introduction of a security control reduce the attack surface in your environment?*





# Democratising Offensive Security

Attack is hard!

Table stakes for modern, cloud offensive security assessment:

- In depth K8s security knowledge
- In depth OS security knowledge
- The “*evil bit*”
- Sufficient time



# Democratising Offensive Security

## Attacker mindset helps defenders

Generally accepted that some form of offensive mindset helps defenders.  
BUT:

- Attack is hard!
- Defenders have a day job
- Offense evolves fast





# Scale

## Datadog case study

The Datadog environment is **VAST** - exactly how vast is confidential 🙈

- *“tens of thousands of nodes”*
- *“hundreds of thousands of pods”*
- *“multi-cloud”*

Traditional **penetration testing does not scale** to this level



# The Solution

Graph theory X offensive security




# Graph Theory

*The study of graphs, mathematical structures used to model pairwise relations between objects.*



# Graph Theory

*The study of graphs, mathematical structures used to model pairwise relations between objects.*





# Graph Theory 101

## Graph

A data type to represent complex, non-linear relationships between objects

- In KubeHound: a Kubernetes cluster

## Vertex

The fundamental unit of which graphs are formed (also known as "node")

- In KubeHound: containers, pods, endpoints

## Edge

A connection between vertices (also known as "relationship")

- In KubeHound: a container escape (e.g *CE\_MODULE\_LOAD*) connects a container and a node

## Path

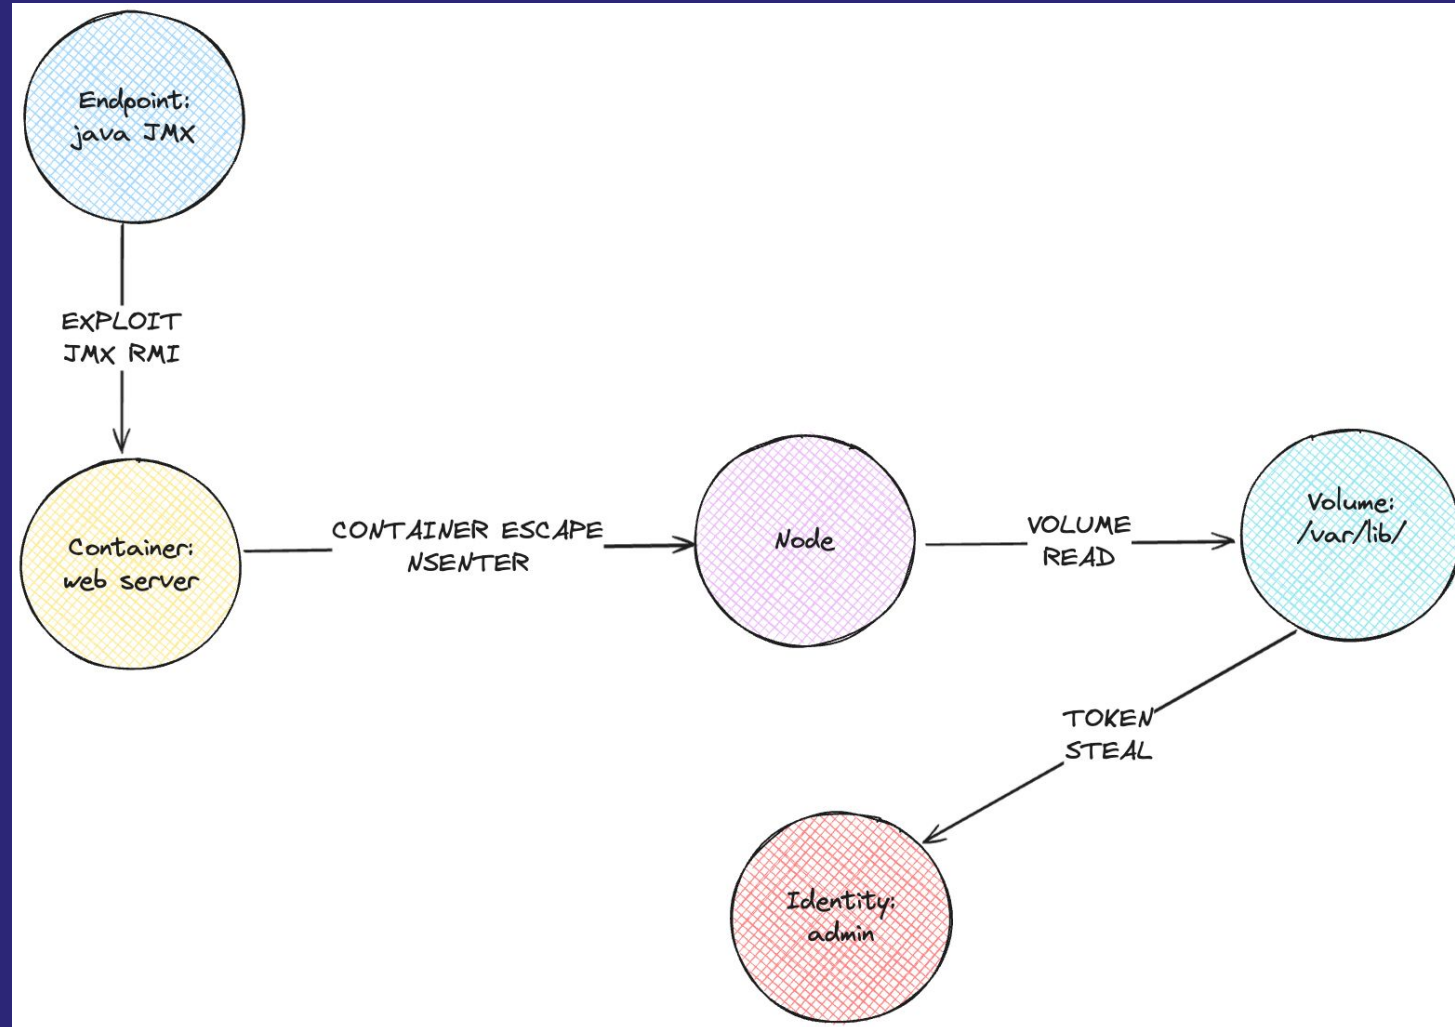
A sequence of edges which joins a sequence of vertices

- In KubeHound: a sequence of attacks from a service endpoint to a cluster admin token



# Attack Graphs

## Attacker TTPs as a data structure



# KubeHound 101

## Entity

An abstract representation of a Kubernetes component that form the **vertices** of the graph

## Attacks

All **edges** in the KubeHound graph represent attacks with a net "improvement" in an attacker's position or a lateral movement opportunity

## Critical Asset

An entity in KubeHound whose compromise would result in cluster **admin** (or equivalent) level access

## Critical Path

A set of connected vertices in the graph that **terminates at a critical asset**





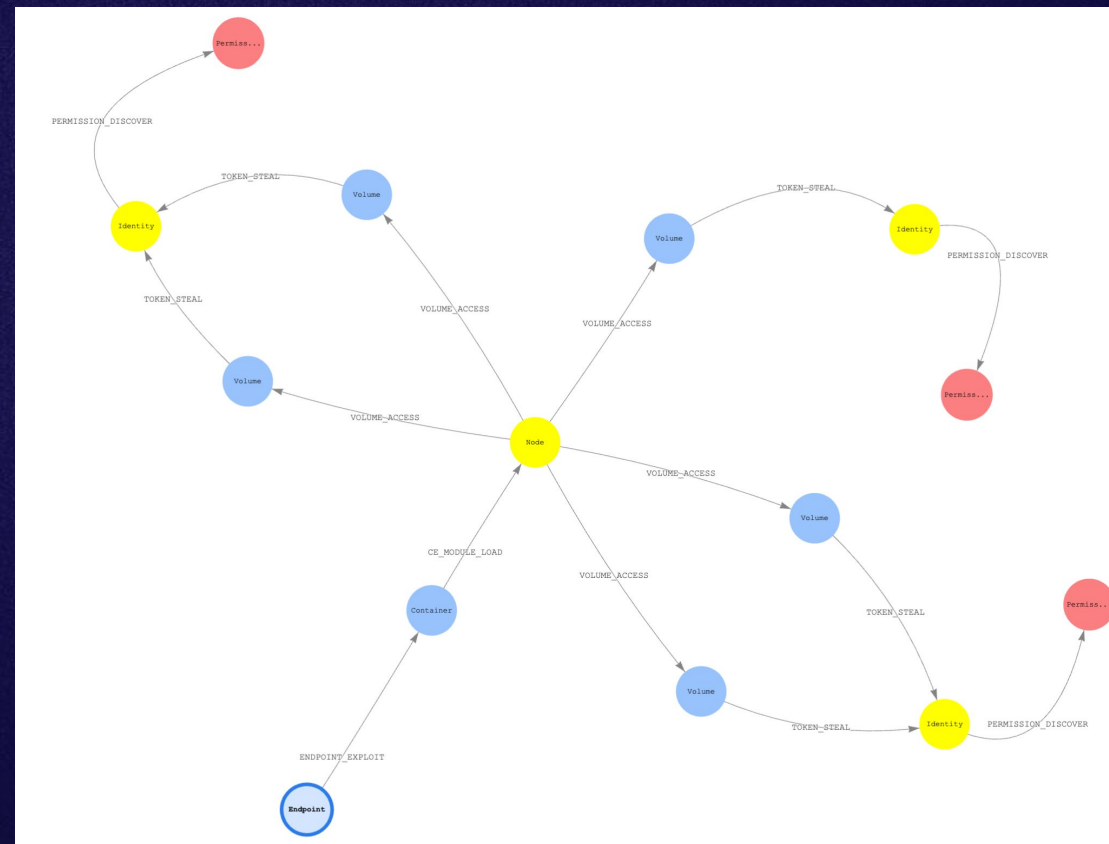
# KubeHound

## Attack Graph

KubeHound creates a graph of attack paths in a Kubernetes cluster, allowing you to identify direct and multi-hop routes an attacker is able to take, visually or through graph queries

## Runtime Calculation

If **any** entity is **connected** to a critical asset in our attack graph - a compromise results in complete control of the cluster





# KubeHound in Action

Case study and capability showcase

# Case Study I

## Modelling a complex attack chain

Can we use KubeHound to find a very complex path (10 hops) from an exposed service to a critical asset?





# Demo

10 hops to cluster admin



# Case Study I

## Modelling a complex attack chain

- ✓ Vulnerability context
- ✓ Scale
- ✓ Democratising offense



# Case Study II

## Quantitative analysis of security posture

Can we use KubeHound to answer the question of “*how secure is my cluster?*” and track that metric over time?





# Demo

Security metrics calculation



# Case Study II

## Quantitative analysis of security posture

- ✓ Quantifying security posture
- ✓ Scale
- ✓ Democratising offense



# Development Process

Research, abstract, design, iterate



# Showcase development process

**Powerful approach - can be generalised**

The approach outlined could be applied to create attack graphs in any number of systems

- Workstations / endpoints
- SaaS identity
- HashiCorp Vault
- AWS



# Research

Domain-specific graph model



# Research I

## Gr0k K8s security

Collate, ingest and categorize all the Kubernetes security research of the past 5 years:

- Conference talks
- Blogs
- Tooling
- Videos



# Research II

## Sketch attack components

Example attack: *“Compromising Kubernetes Cluster by Exploiting RBAC Permissions”* - CyberArk @ RSAC 2020

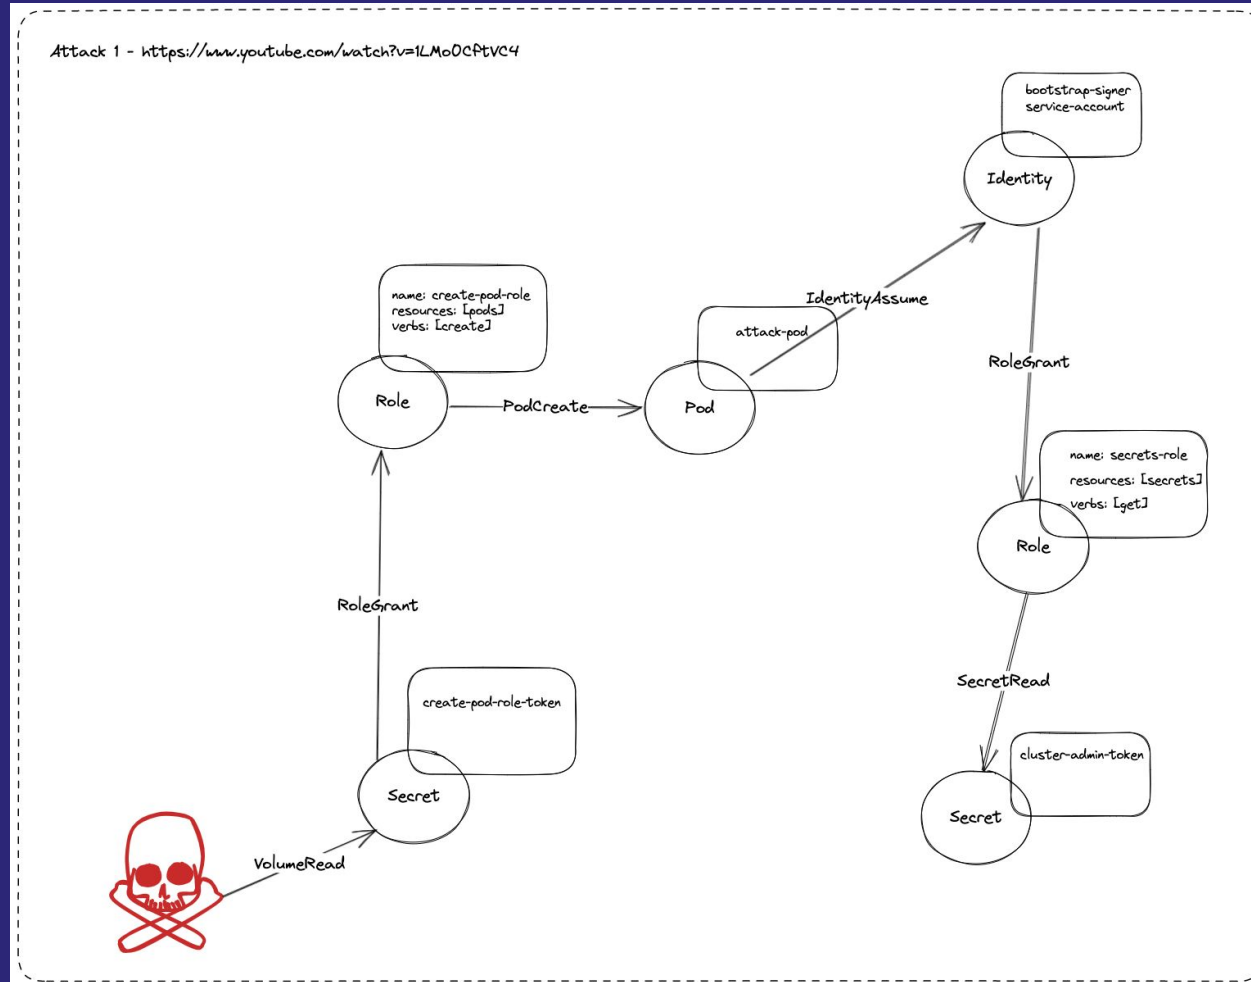
- Attacker compromises a token with **pod/create** permission
- Attacker creates a custom container with an entrypoint script
- Script requests all cluster secrets from K8s API and POSTs to attacker controlled endpoint
- Attack creates a pod with a privileged identity (bootstrap-signer) to run the container image
- *“All your secrets are belong to us”*





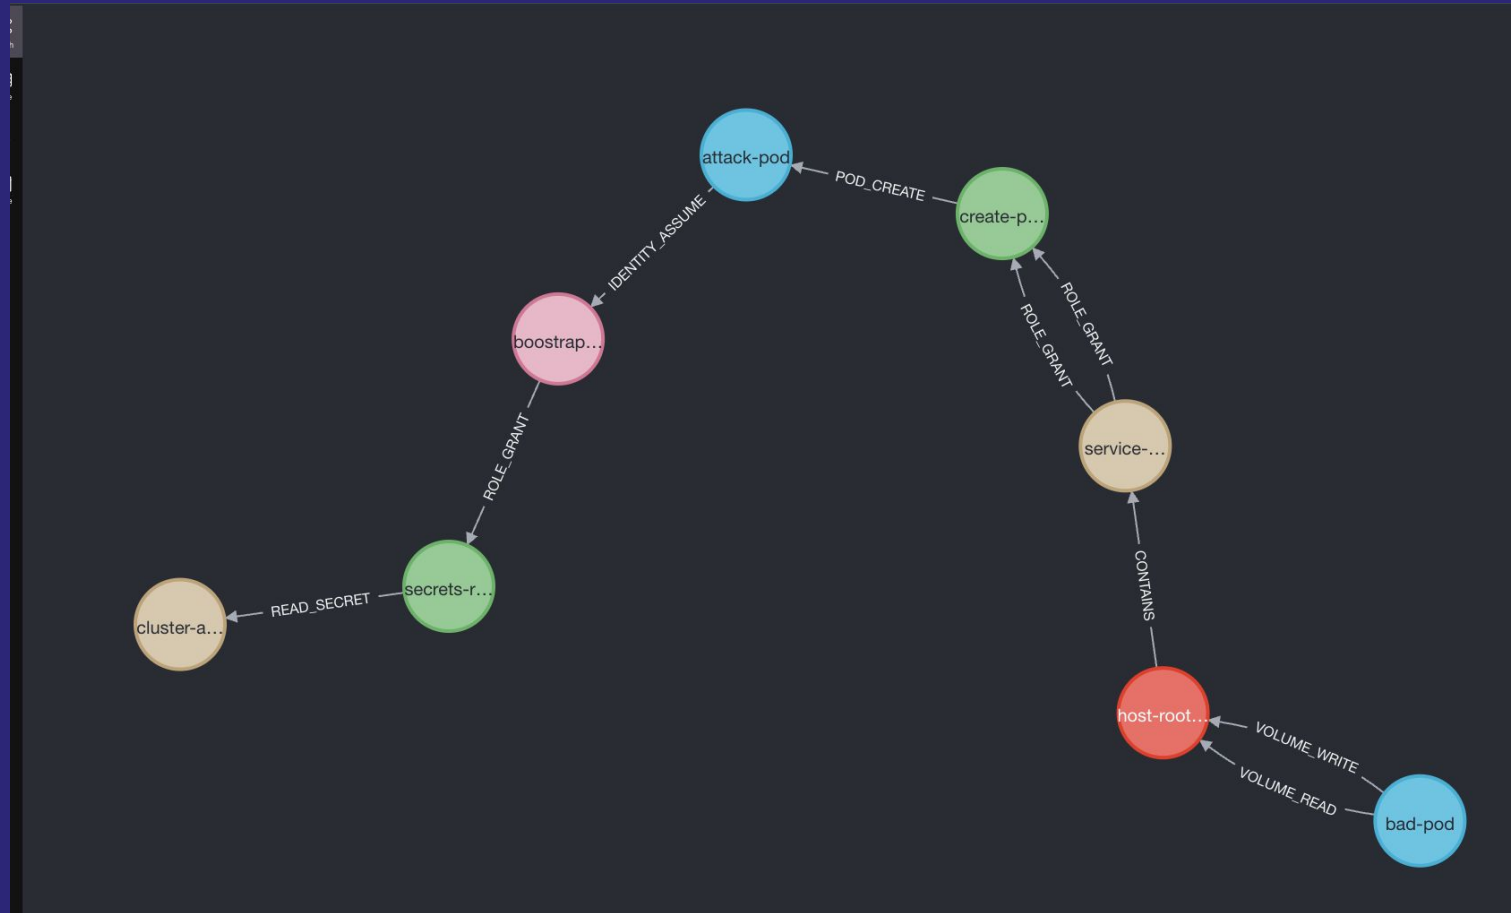
# Research II

## Sketch attack components



# Research III

## Port to graph database



# Research IV

Repeat for all attacks

CE\_MODULE\_LOAD  
CE\_NSENTER  
CE\_PRIV\_MOUNT  
CE\_SYS\_PTRACE  
CE\_UMH\_CORE\_PATTERN  
CE\_VAR\_LOG\_SYMLINK  
CONTAINER\_ATTACH  
ENDPOINT\_EXPLOIT  
SHARE\_PS\_NAMESPACE

EXPLOIT\_CONTAINERD  
EXPLOIT\_HOST\_READ  
EXPLOIT\_HOST\_TRAV  
EXPLOIT\_HOST\_WRITE  
TOKEN\_BRUTEFORCE  
TOKEN\_LIST  
TOKEN\_STEAL  
VOLUME\_ACCESS  
VOLUME\_DISCOVER

IDENTITY\_ASSUME  
IDENTITY\_IMPERSONATE  
PERMISSION\_DISCOVER  
POD\_ATTACH  
POD\_CREATE  
POD\_EXEC  
POD\_PATCH  
...



# Research Phase Outcomes

## Confidence in a notional graph model

Mapping known attacks to a notional model gives us confidence that:

- Our model is complete
- We are not processing unnecessary data

Doing this early, and **thoroughly** ensures we have the correct abstraction





# Design

Efficient translation of source data to graph



# Design I

## Formalise graph model

### Vertices

- Label (i.e entity type)
- Properties (name, namespace, etc...)

### Edges

- Label (i.e attack type)
- Properties (weight)

### Calculation methodology

- Source data
- Algorithm



# Design II

## Data pipeline

### Source ingest

- Now we know what data we need - how do we get it?

### Vertex creation

- What Kubernetes resources do we need to ingest (*pods, nodes, endpoints, etc*)?

### Edge creation

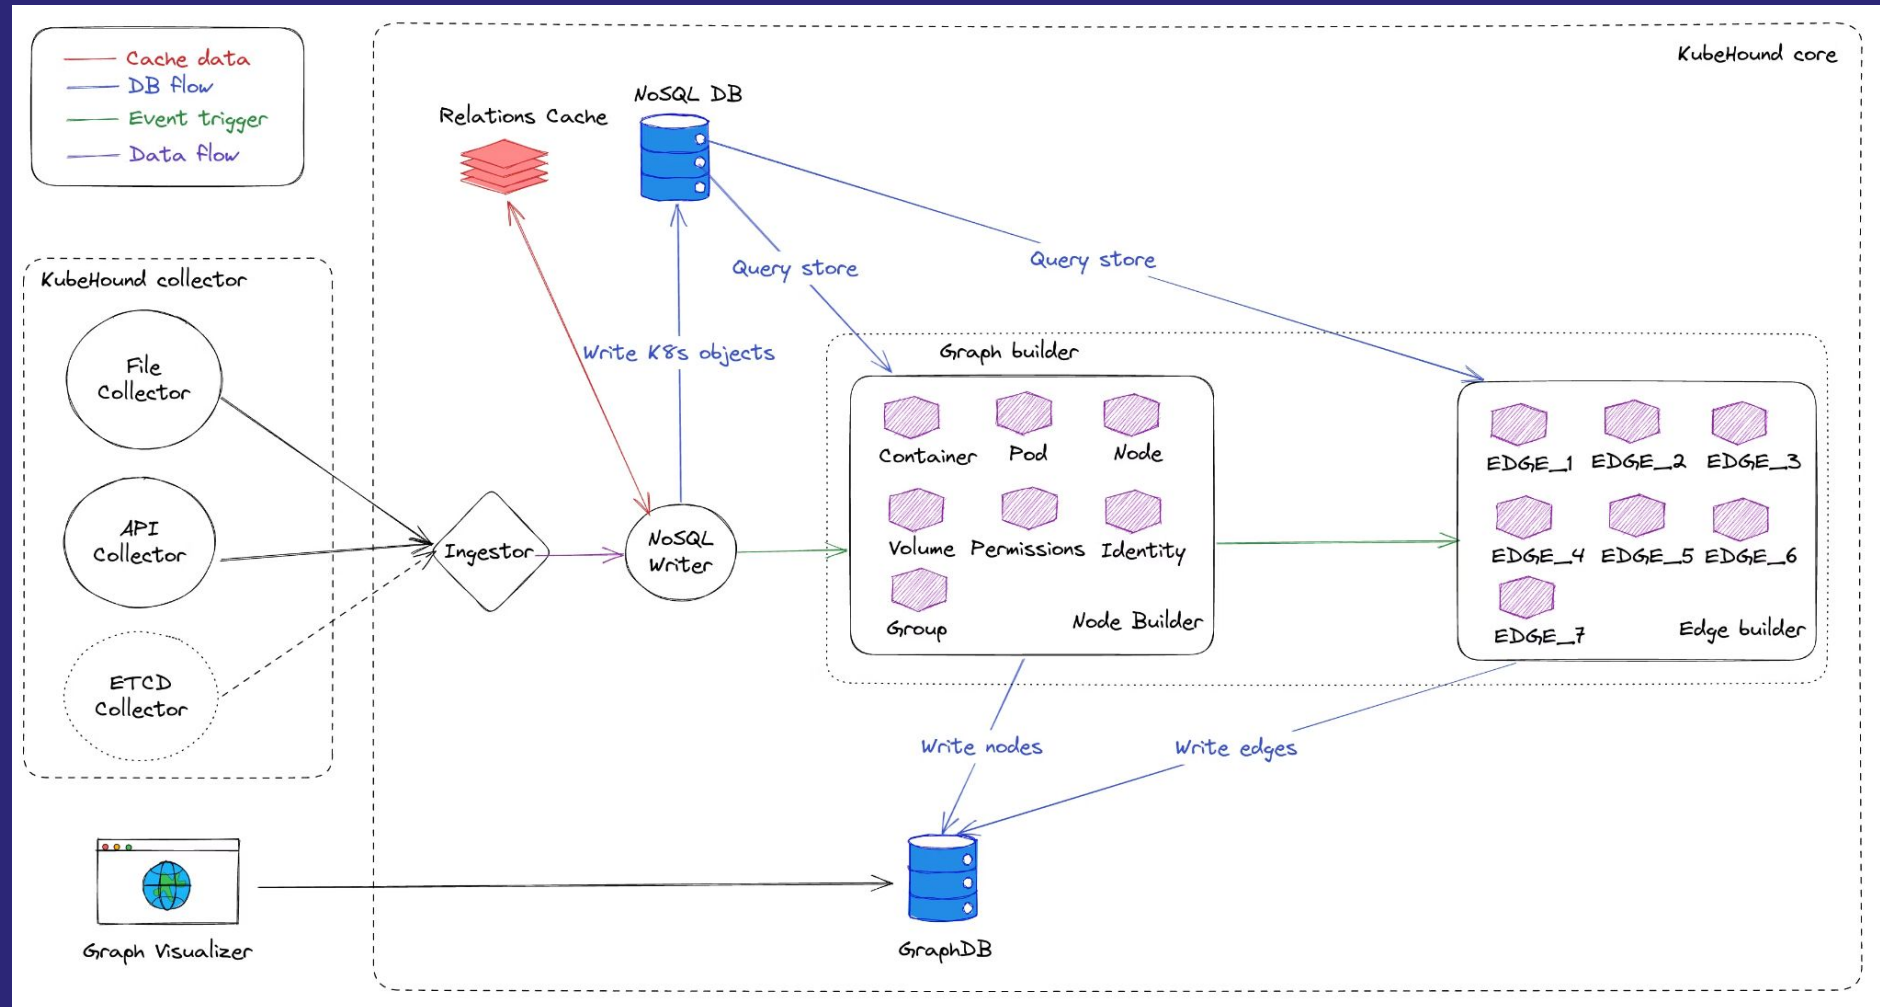
- Do we have the required data to support our edge algorithms?

### Dependencies

- What depends on what?
- What can be done in parallel?
- What must be done in sequence?



# Final Architecture





# Implementation

Guiding principles and outcomes



# KubeHound Development

## Key principles

- Build for scale
- Focus on performance
- Everything will change
- Build for OSS release
- Optimize for ease of contribution of new attacks



# KubeHound Development

## Iterative performance improvement

Notional cluster of 25,000 pods:

- v0.1: 10 hours
- v1.0: 1 hour
- v1.1: 5 minutes\*

\*2 minutes without API rate limiting



# Conclusion

Graphs change the game



# Conclusion

## Automated attack graphs change the game

Some insights gleaned on the power of automated attack graphs through developing and using KubeHound:

- Provide the ability to quantify security posture and risk
- Scale horizontally to handle any environment
- Act as a force multiplier by sharing the mindset of the best offensive practitioners with defenders

**TLDR:** Attack graphs change the game and will be the natural evolution of security tooling





# Thank you

Visit [kubehound.io](https://kubehound.io) to try for yourselves!

Jeremy Fox | [@0xff6a](https://twitter.com/0xff6a)

