

# ADCS: A Love Story

Not really

# Rohan Vazarkar

@CptJesus

Staff Software Developer  
BloodHound Enterprise Team



# The Problem

“When is BloodHound going to support ADCS?”

“ADCS when?”

“Is BloodHound going to get support for ADCS?”

“Can we get an estimate on when ADCS will get implemented???”

# Initial Refinement

- Jonas Knudsen (@jonasbk) and Andy Robbins (@\_wald0) spent a lot of time before engineering work started parsing ADCS research into what the graph should look like
- Determined which escalations would fit in our graph, as well as the intermediary steps necessary for every single escalation
- Determined what data would need to be collected from AD to make sure we had everything accurate

# Initial Estimates

- Sitting in the Seattle office with Jonas, talking about what's necessary for ADCS, thought it wouldn't be so bad
- Filled with hopeful optimism
- Started work, made some progress, estimated that we would be done Q4 of last year
- Started talking about ADCS publicly

# Crushing Defeat

- Turns out, we're (honestly, everyone is) really bad at estimating
- ADCS ended up being about 50 times more complex than we anticipated
- Some of our fundamental assumptions ended up being wrong
- Our traversal engine didn't even support some of the things necessary to do ADCS post processing
- Initial prototypes of ADCS were blowing out RAM on fairly small test environments
- Utter complete sadness

# ADCS Complexity

- ADCS post processing is the most complex logic we've ever put into BloodHound
- Previous “complex” edges like DCSync/SyncLapsPassword were a combination of 2 privileges
- ADCS edges are combinations of 3-5 permission sets

MATCH p1 =  
(n:Base)-[:Enroll|GenericAll|AllExtendedRights|MemberOf\*1..]->(ct:CertTemplate)-[:PublishedTo]-(ca:EnterpriseCA)-[:IssuedSignedBy\*1..]-(rootca:RootCA)-[:RootCAFor]->(d:Domain)

MATCH p2 = (n)-[:Enroll|GenericAll|AllExtendedRights|MemberOf\*1..]->(ca)-[:TrustedForNTAuth]->(nt:NTAuthStore)-[:NTAuthStoreFor]-(d)

# The Easy Way



- Find all the first degree controllers of node A and node B
- Expand all group members recursively
- Find all the intersections between the two
- Make an edge from each principal found this way as an attack path



# The hard, more right way

- Use the compressed bitmap to store all our group membership expansions to reduce expensive traversal time
- Implement shortcutting logic to opportunistically create edges at the highest possible level, minimizing the number of created edges and taking advantage of graph semantics
  - If the Domain Users group has the permission, all the members of the group have the permission via the MemberOf edge in the graph
  - Cost - some accuracy is lost as some ancillary, valid attack paths can be obscured by the shortcutting

# Shortcutting - These graphs mean the same thing!



# Edge Composition

- Post processing is hard
- Performantly evaluating the entire dataset to make sure that we catch edge cases takes effort, but we're working with no start/end point so we can just batch up logic
- Edge composition is even harder as its restricted to a specific set of nodes instead of just being able to use large sets of nodes

```

MATCH (u:User {objectid:'S-1-5-21-2057499049-1289676208-1959431660-238209'})-[:ADCSesC1]->(d:Domain {objectid:'S-1-5-21-1621856376-872934182-3936853371'})
MATCH (c:Container)-[:Contains]->(rca:RootCA)
WHERE c.name CONTAINS "CERTIFICATION AUTHORITIES" AND c.domain = d.domain
MATCH (ca:EnterpriseCA)-[:IssuedSignedBy|EnterpriseCAFor*1..]->(rca)
WHERE (ca)-[:TrustedForNTAuth]->(:NTAuthStore)
MATCH (ct:CertTemplate)-[:PublishedTo]->(ca)
WHERE (ct.requiresmanagerapproval = false
AND ct.schemaversion > 1
AND ct.authorizedsignatures = 0
AND ct.authenticationenabled = true
AND ct.enrolleesuppliessubject = true)
OR (ct.requiresmanagerapproval = false
AND ct.schemaversion = 1
AND ct.authenticationenabled = true
AND ct.enrolleesuppliessubject = true)
OPTIONAL MATCH p1 = (u)-[:GenericAll|Enroll|AllExtendedRights]->(ct)-[:PublishedTo]->(ca)
OPTIONAL MATCH p2 = (u)-[:MemberOf*1..]->(:Group)-[:GenericAll|Enroll|AllExtendedRights]->(ct)-[:PublishedTo]->(ca)
OPTIONAL MATCH p3 = (u)-[:Enroll]->(ca)-[:IssuedSignedBy|EnterpriseCAFor|RootCAFor*1..]->(d)
OPTIONAL MATCH p4 = (u)-[:MemberOf*1..]->(:Group)-[:Enroll]->(ca)-[:IssuedSignedBy|EnterpriseCAFor|RootCAFor*1..]->(d)
OPTIONAL MATCH p5 = (ca)-[:TrustedForNTAuth]->(:NTAuthStore)-[:NTAuthStoreFor]->(d)
RETURN p1,p2,p3,p4,p5

```

# Fix it!

- Our traversal engine for cypher was missing several pieces necessary for these
- Spent 2 weeks just rewriting our engine to allow us to begin doing edge composition - pattern continuations allow us to write GO code that mimics the cypher previously shown in a very logical way

```

func adcsESC13Path1Pattern() traversal.PatternContinuation {
    return traversal.NewPattern().
        OutboundWithDepth(
            min: 0, max: 0,
            query.And(
                query.Kind(query.Relationship(), ad.MemberOf),
                query.Kind(query.End(), ad.Group),
            ),
        ).
        OutboundWithDepth(
            min: 1, max: 1,
            query.And(
                query.KindIn(query.Relationship(), ad.Enroll, ad.GenericAll, ad.AllExtendedRights),
                query.KindIn(query.End(), ad.CertTemplate),
                query.And(
                    query.Equals(query.EndProperty(ad.AuthenticationEnabled.String()), value: true),
                    query.Or(
                        query.Equals(query.EndProperty(ad.SchemaVersion.String()), value: 1),
                        query.And(
                            query.GreaterThan(query.EndProperty(ad.SchemaVersion.String()), value: 1),
                            query.Equals(query.EndProperty(ad.AuthorizedSignatures.String()), value: 0),
                        ),
                    ),
                    query.Equals(query.EndProperty(ad.RequiresManagerApproval.String()), value: false),
                ),
            ),
        ).OutboundWithDepth(
            min: 1, max: 1,
            query.And(
                query.KindIn(query.Relationship(), ad.PublishedTo),
                query.Kind(query.End(), ad.EnterpriseCA),
            ),
            Outbound(query.And(
                query.KindIn(query.Relationship(), ad.IssuedSignedBy, ad.EnterpriseCAFor),
                query.Kind(query.End(), ad.RootCA),
            )),
            Outbound(query.And(
                query.KindIn(query.Relationship(), ad.RootCAFor),
                query.KindIn(query.End(), ad.Domain),
            ))
        )
}

```

# Did it work?

- Wrote the first edge composition POC - immediately killed our test instance by blowing out memory
- Rewrote our edge composition completely using incredibly gross optimizations - finally it worked!
  - Opportunistic pruning of target nodes
  - Starting from the middle of a query and working backwards to reduce total path space

# The problem?

- We're now late Q4 of 2023.
- We're just wrapping up ESC1. We need to add ESC3,4,6,9,10,13
- Jonas and Andy are still finding more escalations
- Clearly our timelines were correct
- Finally just accepted that we couldn't deliver Q4, doubled down on the work and got it all done in Q1 of 2024 with extra engineering effort



# Take-aways

- Complex problems are extremely difficult to estimate timelines for
- Some problems are vastly more complex than they seem on the surface
- Engineers think they can solve everything quickly and perfectly (especially me)
- I'm just rambling at this point